# 3.5 Advanced Feature Extraction (cont'd)

```
In [1]: # Library Imports:
        import warnings
        warnings.filterwarnings('ignore')
        import numpy as np
        import pandas as pd
        import matplotlib.pyplot as plt
        %matplotlib inline
        import seaborn as sns
        from subprocess import check_output
        import os
        import gc
        from bs4 import BeautifulSoup
        from os import path
        from PIL import Image
        # Importing the foll. library show a word-cloud plot:
        # https://stackoverflow.com/questions/45625434/how-to-install-wordcloud-in-python3-6
        from wordcloud import WordCloud, STOPWORDS
        import time
```

# 3.5.1 Analysis of Extracted Features

#### 3.5.1.1 Plotting Word Clouds

- Creating word clouds of Duplicate Questions and Non-duplicate Questions.
- We can observe which of the words occur most frequently.

In [3]: # Loading the preprocessed & advanced featurized data:
 df = pd.read\_csv('nlp\_features\_train.csv', encoding='latin-1')
 print(df.shape)
 df.head()

(404290, 21)

Out[3]:

: [	Τ			48 6		l														_
	id	qid1	qid2	question1	question2	is_duplicate	cwc_min	cwc_max	csc_min	csc_max		ctc_max	last_word_eq	first_word_eq	abs_len_diff	mean_len	token_set_ratio	token_sort_ratio	fuzz_ratio	fu
	0	1	2	step by step guide to invest in	13	0	0.999980	0.833319	0.999983	0.999983	::	0.785709	0.0	1.0	0.0	13.0	100	93	93	10
	1 1	3	4	story of kohinoor koh i noor	what would happen if the indian government sto		0.799984	0.399996	0.749981	0.599988		0.466664	0.0	1.0	0.0	12.5	86	63	66	75
	2	5	6	increase the speed of my internet	how can internet speed be increased by hacking	0	0.399992	0.333328	0.399992	0.249997		0.285712	0.0	1.0	0.0	12.0	63	63	43	47
	3 3	7	8	mentally very lonely how can i	find the remainder when math 23 24 math i	0	0.000000	0.000000	0.000000	0.000000		0.000000	0.0	0.0	0.0	12.0	28	24	9	14
	4	9	10	which one dissolve in water quikly sugar salt	would	0	0.399992	0.199998	0.999950	0.666644		0.307690	0.0	1.0	0.0	10.0	67	47	35	56

5 rows × 21 columns

```
In [6]: df dup = df[df.is duplicate == 1]
         df non dup = df[df.is duplicate == 0]
         # Converting 2D array of q1 & q2 into 1D array by flattening the array:
         # Flattening Example: [[1,2], [3,4]] ---flatten()---> [1,2,3,4]
         dup = np.dstack([df_dup.question1, df_dup.question2]).flatten()
         non dup = np.dstack([df non dup.question1, df non dup.question2]).flatten()
         print("Number of data points in class 1 [duplicate question pairs]:", len(dup))
         print("Number of data points in class 0 [non-duplicate question pairs]:", len(non_dup))
         # Saving the numpy arrays dup and non_dup into a text file:
         np.savetxt('train_dup_question_pairs.txt', dup, delimiter=' ', fmt='%s')
         np.savetxt('train non dup question pairs.txt', non dup, delimiter=' ', fmt='%s')
         Number of data points in class 1 [duplicate question pairs]: 298526
         Number of data points in class 0 [non-duplicate question pairs]: 510054
In [10]: # Reading the text files and removing the Stop Words:
         d = path.dirname('.')
         text dup w = open(path.join(d, 'train dup question pairs.txt')).read()
         text_non_dup_w = open(path.join(d, 'train_non_dup_question_pairs.txt')).read()
         stopwords = set(STOPWORDS)
         stopwords.add('said')
         stopwords.add('br')
         stopwords.add(' ')
         stopwords.remove('not')
         stopwords.remove('no')
         stopwords.remove('like')
         print("Total number of words in duplicate pair questions:", len(text_dup_w))
         print("Total number of words in non duplicate pair questions:", len(text_non_dup_w))
         Total number of words in duplicate pair questions: 16100349
```

Word Clouds generated from Duplicate Pair Question's text:

Total number of words in non duplicate pair questions: 33168037

```
In [12]: wc = WordCloud(background_color="white", max_words=len(text_dup_w), stopwords=stopwords)
    wc.generate(text_dup_w)
    print('Word Cloud for Duplicate Question Pairs:')
    plt.imshow(wc, interpolation='bilinear')
    plt.axis('off')
    plt.show()
```

Word Cloud for Duplicate Question Pairs:



The word cloud above shows that the most frequently occurring words in all the duplicate questions are 'best way', 'donald trump', '1k rupee', etc. The bigger the word in the word cloud, the more frequently occurring the word is.

#### Word Cloud generated from Non-duplicate Pair Question's text:

```
In [13]: wc1 = WordCloud(background_color="white", max_words=len(text_non_dup_w), stopwords=stopwords)
    wc1.generate(text_non_dup_w)
    print('Word Cloud for Non-Duplicate Question Pairs:')
    plt.imshow(wc1, interpolation='bilinear')
    plt.axis('off')
    plt.show()
```

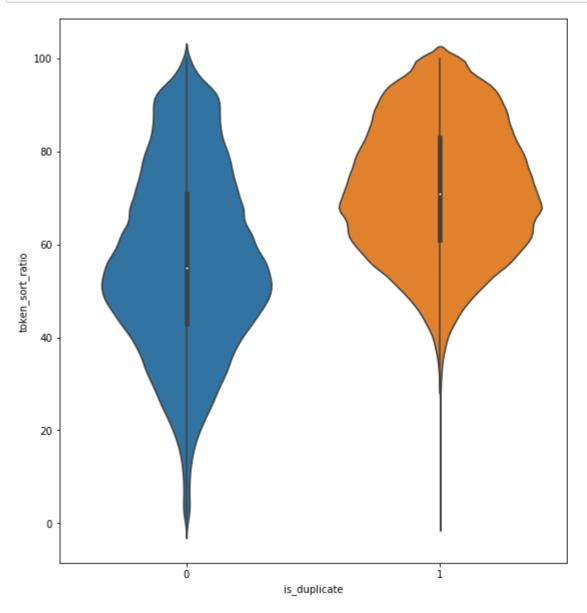
Word Cloud for Non-Duplicate Question Pairs:

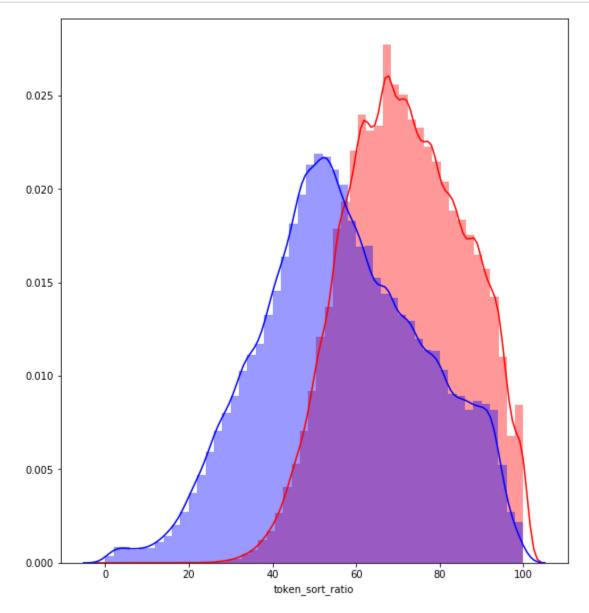


The word cloud above shows that the most frequently occurring words in Non-Duplicate Question Pairs are 'not', 'difference', 'india', etc.

```
In [14]: sns.pairplot(
                df[['ctc_min', 'cwc_min', 'csc_min', 'token_sort_ratio', 'is_duplicate']]\
                [0:df.shape[0]], hue = 'is_duplicate',
                vars=['ctc_min', 'cwc_min', 'csc_min', 'token_sort_ratio']
           plt.show()
               1.0
                                           1.0 -
                                                                                                 1.0
                0.8
                                           0.8
             0.6
E 0.4
                                           0.6
                                                                      0.6
                                           0.4
                                                                      0.4
                0.2
                                           0.2
                                                                      0.2
                0.0
                                           0.0
                                                                      0.0
                                                                1.0
                                                                                                                      100
                    0.0
                            0.5
                                    1.0
                                              0.0
                                                       0.5
                                                                          0.0
                                                                                           1.0
                1.0
                                           1.0
                                                                      1.0 -
                                                                                                 1.0
                0.8
                                           0.8
                                                                                                 0.8
             0.6 ·
                                           0.6
                                                                      0.6
                                           0.4
                0.2
                0.0
                                           0.0
                                                                      0.0
                                                                                                 0.0
                                    1.0
                                                                                           1.0
                                              0.0
                                                                                                                      100
                                                                                                                          is duplicate
                1.0
                                           1.0 -
                                                                      1.0
                                                                                                 1.0 -
                                                                                                                             • 1
                0.8
                                           0.8
                                                                      0.8
                                                                                                 0.8
             .드 0.6 ·
당 0.4 ·
                                                                      0.6
                                                                                                 0.6
                                           0.6
                                                                                                 0.4
                0.2
                                                                                                 0.2
                                           0.2
                0.0
                                                                                                 0.0
                                           01/400 (0 01/4000:0 (0 f/0 (00 0)))
                                    1.0
                                                                1.0
                    0.0
                                                                                                                      100
                                               0.0
                                                       0.5
                                                                          0.0
                                                                                           1.0
               100
                                                                                                100
                                            80
                                                                      80
                                                                                                  80
             token sort ratio
                60
                                                                                                  60
                                                                                                  40
                20
                                                                                                  20
                                                                                           1.0
                    0.0
                            0.5
                                     1.0
                                               0.0
                                                       0.5
                                                                1.0
                                                                          0.0
                                                                                  0.5
                                                                                                              50
                                                                                                                      100
                          ctc_min
                                                     cwc_min
                                                                                 csc_min
                                                                                                         token_sort_ratio
```

Analysis: There is signficant overlap in all the pairplots, but we can still see that there are blue and orange points visible separately.





If we see both the plots above, with the PDF distribution, we can see that there's an overlap to near right of the blue PDF with the red PDF.

With the Violin plots, we can see that the 50th percentiles (medians) of both duplicate & non-duplicate quesion pairs are not overlapping.

Therefore, this feature 'token\_sort\_ratio', plays an important role for our duplicacy check.


```
In [17]: # Distribution of the fuzz_ratio:
    plt.figure(figsize=(20,10))


# Violin Plot for fuzz_ratio feature:
    plt.subplot(1,2,1)
    sns.violinplot(x = 'is_duplicate', y = 'fuzz_ratio', data = df[0:])

# Histogram / PDF of fuzz_ratio feature:
    plt.subplot(1,2,2)
    sns.distplot(
        df[df.is_duplicate == 1.0]['fuzz_ratio'][0:],
        label='1', color='red'
)

sns.distplot(
    df[df.is_duplicate == 0.0]['fuzz_ratio'][0:],
    label='0', color='blue'
)

plt.show()
```





Even the fuzz ratio PDF plot has significant overlap in the middle. But we can clearly see some red points on the top right, and blue points on the left side of the PDF plot.

With the Violin Plots, we can see that the 50th percentiles (medians) of both duplicate & non-duplicate question pairs are not overlapping. That is a good sign, that this feature is highly useful, when we apply an ML model to classify whether the questions are duplicate or not.

# 3.5.2 Visualization using PCA & t-SNE

We are now visualizing only advanced features. Later we will visualize 2D plot of all the features along with word2vec vector representations of q1 & q2.

#### 3.5.2.1 Visualization using Principal Component Analysis [PCA]

~> The shape of the sampled data is: (10000, 20)
~> The shape of the class label is: (10000,)

We will first see a 2D plot, and then see a 3D plot.

#### **2D PCA Plot**

```
In [30]: # We will apply PCA on sample data (10k data points):
    df_s = df[0:10000]

# Sample the class Label (is_duplicate) from sampled data into a new variable:
    sample_labels = df_s['is_duplicate']

# Drop the class Label (is_duplicate) from the sampled data:
    df_sample = df_s.drop('is_duplicate', axis=1)

print("~> The shape of the sampled data is: {}".format(df_sample.shape))
    print("~> The shape of the class label is: {}".format(sample_labels.shape))
```

```
Out[31]:
            id | qid1 | qid2 | question1 |
                                  question2 | cwc_min | cwc_max | csc_min | csc_max | ctc_min | ctc_max | last_word_eq | first_word_eq | abs_len_diff | mean_len | token_set_ratio | token_sort_ratio | fuzz_ratio | fuzz_par
                        what is the what is the
         0
                                 step by
                        step by
                                           0.999980 | 0.833319 | 0.999983 | 0.999983 | 0.916659 | 0.785709 | 0.0
                                                                                                            1.0
                                                                                                                        0.0
                                                                                                                                   13.0
                                                                                                                                             100
                                                                                                                                                          93
                                                                                                                                                                         93
                                                                                                                                                                                   100
            0 1
                        step guide step guide
                        to invest in to invest in
                                 sh...
                        what is the what would
                        story of
                                 happen if
            1 3
                                 the indian
                                           0.799984 | 0.399996 | 0.749981 | 0.599988 | 0.699993 | 0.466664 | 0.0
                                                                                                            1.0
                                                                                                                        0.0
                                                                                                                                   12.5
                                                                                                                                             86
                                                                                                                                                          63
                                                                                                                                                                         66
                                                                                                                                                                                  75
                        kohinoor
                                 government
                        koh i noor
                        dia...
                                 sto...
In [32]: sample_labels.head()
Out[32]: 0
              0
         1
         2
              0
         3
              0
         4
              0
         Name: is duplicate, dtype: int64
In [37]: # Data-preprocessing: Standardize the data, i.e., Mean Centering:
         from sklearn.preprocessing import StandardScaler
         # We will not consider q1 & q2, because they're not converted into numeric vector till now
         std data = StandardScaler().fit transform(
             df_sample[['cwc_min', 'cwc_max', 'csc_min', 'csc_max',
                        'ctc_min', 'ctc_max', 'last_word_eq', 'first_word_eq',
                        'abs_len_diff', 'mean_len', 'token_set_ratio', 'token_sort_ratio',
                        'fuzz ratio', 'fuzz parital ratio', 'longest substr ratio']]
         print(std_data.shape)
         (10000, 15)
In [39]: # Check the stdized data:
         print(std_data[:2,:])
         -0.6762404 0.97140856 0.
                                              0.33699245 1.37512904 1.61509069
            1.57318235 1.87063632 2.58200464]
          -0.6762404 0.97140856 0.
                                              0.24248297 0.69698109 0.03769451
            0.34693315 0.54733006 0.89940355]]
```


In [31]: df sample.head(2)

```
In [40]: # Library to use PCA:
         from sklearn import decomposition
         # Initiate the PCA instance:
         pca = decomposition.PCA()
         st = time.time()
         # Configuring the parameters:
         # The number of components = 2, since we are projecting 15D points on a 2D plane.
         # We will get the top 2 eigen vectors as the principal components:
         pca.n_components = 2
         pca_sample_data = pca.fit_transform(std_data)
         print("~> Shape of pca_sample_data is:", pca_sample_data.shape)
         print('Total time taken to convert 15D data projection on 2D is:', \
              time.time()-st, 'seconds')
         ~> Shape of pca_sample_data is: (10000, 2)
         Total time taken to convert 15D data projection on 2D is: 0.9166522026062012 seconds
In [41]: # Attaching the sample_label feature to the attained pca_sample_data to plot
         # our data properly:
         pca_sample_data_final = np.vstack((pca_sample_data.T, sample_labels)).T
         pca_sample_data_final.shape
Out[41]: (10000, 3)
In [61]: # Creating a new dataframe which will help us in easier plotting:
         df_pca_final = pd.DataFrame(
                             data=pca_sample_data_final,
                             columns = ('1st_PC', '2nd_PC', 'is_duplicate')
         df_pca_final.head()
```

Out[61]:

	1st_PC	2nd_PC	is_duplicate
0	-5.152321	0.759183	0.0
1	-1.427434	0.786293	0.0
2	1.740200	0.603417	0.0
3	6.480144	-0.076563	0.0
4	0.948609	1.956853	0.0


```
In [70]: # Plot both 1st_PC and 2nd_PC using Seaborn, where hue is 'is_duplicate':
    sns.lmplot(
        data=df_pca_final, x='1st_PC', y='2nd_PC', hue='is_duplicate',
        fit_reg=False, size=8, palette="Set1", markers=['s','o']
    )
    plt.title('15 Dimensionsional Data Projected on 2 Prinicipal Components')
    plt.show()
```



From the PCA plot, we can see that when we try to visualize our data in 2D, we have overlap of blue and red points, which we can see evidently from the PCA plot above. There's also one outlier point which is a red point at the top right of the plot.

We will now see the percentage of variance explained at different number of dimensions

```
In [72]: # Generate 15 Dimensional Principal Components:
         pca.n_components = 15
         # Project the 15 Dimensional data on a 15 Dimensional Hyperplane:
         pca_variance = pca.fit_transform(std_data)
         # We will generate individual eigen values divided with the eigen sums:
         # percentage_variance_explained is named as 'perc_var_exp'
         perc_var_exp = pca.explained_variance_ / np.sum(pca.explained_variance_)
         # Now we will take a Prefix Sum of the perc_var_exp:
         ps_var_explained = np.cumsum(perc_var_exp)
         # Plot the prefix summed variances to find the number of dimensions to be taken
         # to retain the information/variance at different amounts of percentages:
         plt.figure(1, figsize=(20,10))
         plt.clf()
         plt.plot(ps_var_explained, linewidth=3)
         plt.grid()
         plt.xlabel('Number of Components ---->')
         plt.ylabel('Cummulative Explained Variance ---->')
         plt.show()
```



Analysis: We can see that more than 95% of variance is explained when we reduce the dimensions of the original 15 dimensional data to 6 dimensional data. With 2 Dimensions, we can only explain more than 80% of data. That's the reason, we can see a significant overlap of red & blue points in the PCA plot.

With reduction to 3 Dimensions, we can explain more than 85% of data.

### **3D PCA Plot**

```
In [73]: # Intitate the PCA instance:
pca_3d = decomposition.PCA()

# Configuring the parameters:
# The number of components = 3, since we are projecting 15D points on a 3D plane.
# We will get the top 3 eigen vectors as the principal components:
pca_3d.n_components = 3

pca_sample_data_3d = pca_3d.fit_transform(std_data)


print("~> Shape of pca_sample_data is:", pca_sample_data_3d.shape)

--> Shape of pca_sample_data is: (10000, 3)

In [74]: # Attaching the sample_label feature to the attained pca_sample_data_3d to plot
# our data properly:
pca_sample_data_final_3d = np.vstack((pca_sample_data_3d.T, sample_labels)).T
pca_sample_data_final_3d.shape
```

Out[74]: (10000, 4)

```
In [100]: # Necessary library to make a 3D Scatter plot:
          %matplotlib notebook
          import matplotlib
          from matplotlib import colors
          from mpl_toolkits.mplot3d import Axes3D
          fig = plt.figure(figsize=(10,8))
          ax = Axes3D(fig)
          color = ['red', 'blue']
          l = [float(i) for i in sample_labels]
          ax.scatter(
              pca_sample_data_final_3d[:,0:1], pca_sample_data_final_3d[:,1:2],
              pca_sample_data_final_3d[:,2:3], c = 1,
              cmap = matplotlib.colors.ListedColormap(color)
          ax.set_title('3D PCA Plot with Engineered Features')
          ax.set_xlabel('1st_Principal_Component')
          ax.set_ylabel('2nd_Principal_Component')
          ax.set_zlabel('3rd_Principal_Component')
          plt.show()
```



Analysis: We can see that there's a lot of overlap even in the 3D PCA plot. Therefore, now we will go for t-SNE plotting, both 2D and 3D.

# 3.5.2.2 Visualization using t-Distributed Stochastic Neighbourhood Embedding [t-SNE]

t-SNE algorithm doesn't depend on a standardized data [mean centered], but depends on normalized data, because here, we are trying t-SNE tries to find distance between points to embed them in a lower dimension, and mean centering won't have any affect when finding the distances.

#### 2D t-SNE Plot

```
In [102]: # Using t-SNE for visualizing 15 Dimensional data in a 2 Dimensional Plane:
# Necessary Library import for Normalizing the data:
from sklearn.preprocessing import MinMaxScaler

# We will use the previously sampled 10k data points, i.e., df_sample
df_sample.head(1)
```

Out[102]:

id	d qio	d1 qid2	question1	question2	cwc_min	cwc_max	csc_min	csc_max	ctc_min	ctc_max	last_word_eq	first_word_eq	abs_len_diff	mean_len	token_set_ratio	token_sort_ratio	fuzz_ratio	fuzz_pari
0	1	2	step by step guide to invest in	what is the step by step guide to invest in sh	0.99998	0.833319	0.999983	0.999983	0.916659	0.785709	0.0	1.0	0.0	13.0	100	93	93	100

As we can see above, all the points are normalized.

0.46666667, 0. , 1. , 0. , 0.09541985,

, 0.75 , 0.60010188]])

, 0.66

, 0.63

```
In [107]: # Necessary Library Imports to apply t-SNE on our data:
          from sklearn.manifold import TSNE
          tsne 2D model = TSNE(
              n components=2, init='random', method='barnes hut', n iter=1000,
              verbose=2, angle=0.5
          st = time.time()
          tsne 2D data = tsne 2D model.fit transform(df sample normalized)
          print('Total time taken to get the tsne 2D data: {} minutes'
                .format(round((time.time()-st)/60, 2)))
          [t-SNE] Computing 91 nearest neighbors...
          [t-SNE] Indexed 10000 samples in 0.086s...
          [t-SNE] Computed neighbors for 10000 samples in 1.387s...
          [t-SNE] Computed conditional probabilities for sample 1000 / 10000
          [t-SNE] Computed conditional probabilities for sample 2000 / 10000
          [t-SNE] Computed conditional probabilities for sample 3000 / 10000
          [t-SNE] Computed conditional probabilities for sample 4000 / 10000
          [t-SNE] Computed conditional probabilities for sample 5000 / 10000
          [t-SNE] Computed conditional probabilities for sample 6000 / 10000
          [t-SNE] Computed conditional probabilities for sample 7000 / 10000
          [t-SNE] Computed conditional probabilities for sample 8000 / 10000
          [t-SNE] Computed conditional probabilities for sample 9000 / 10000
          [t-SNE] Computed conditional probabilities for sample 10000 / 10000
          [t-SNE] Mean sigma: 0.106319
          [t-SNE] Computed conditional probabilities in 0.739s
          [t-SNE] Iteration 50: error = 95.4943619, gradient norm = 0.0220093 (50 iterations in 21.959s)
          [t-SNE] Iteration 100: error = 77.5966644, gradient norm = 0.0044036 (50 iterations in 16.863s)
          [t-SNE] Iteration 150: error = 74.8951645, gradient norm = 0.0026241 (50 iterations in 16.181s)
          [t-SNE] Iteration 200: error = 73.8063049, gradient norm = 0.0017627 (50 iterations in 17.349s)
          [t-SNE] Iteration 250: error = 73.1917343, gradient norm = 0.0015083 (50 iterations in 16.364s)
          [t-SNE] KL divergence after 250 iterations with early exaggeration: 73.191734
          [t-SNE] Iteration 300: error = 2.5369163, gradient norm = 0.0012834 (50 iterations in 17.970s)
          [t-SNE] Iteration 350: error = 2.0162439, gradient norm = 0.0005868 (50 iterations in 18.293s)
          [t-SNE] Iteration 400: error = 1.7574749, gradient norm = 0.0003545 (50 iterations in 17.522s)
          [t-SNE] Iteration 450: error = 1.6041224, gradient norm = 0.0002478 (50 iterations in 17.821s)
          [t-SNE] Iteration 500: error = 1.5031897, gradient norm = 0.0001825 (50 iterations in 17.593s)
          [t-SNE] Iteration 550: error = 1.4315838, gradient norm = 0.0001420 (50 iterations in 17.160s)
          [t-SNE] Iteration 600: error = 1.3783228, gradient norm = 0.0001152 (50 iterations in 17.778s)
          [t-SNE] Iteration 650: error = 1.3374861, gradient norm = 0.0000973 (50 iterations in 17.263s)
          [t-SNE] Iteration 700: error = 1.3056525, gradient norm = 0.0000853 (50 iterations in 17.101s)
          [t-SNE] Iteration 750: error = 1.2804481, gradient norm = 0.0000765 (50 iterations in 17.881s)
          [t-SNE] Iteration 800: error = 1.2608801, gradient norm = 0.0000713 (50 iterations in 17.660s)
          [t-SNE] Iteration 850: error = 1.2457664, gradient norm = 0.0000652 (50 iterations in 18.140s)
          [t-SNE] Iteration 900: error = 1.2331495, gradient norm = 0.0000615 (50 iterations in 17.224s)
```

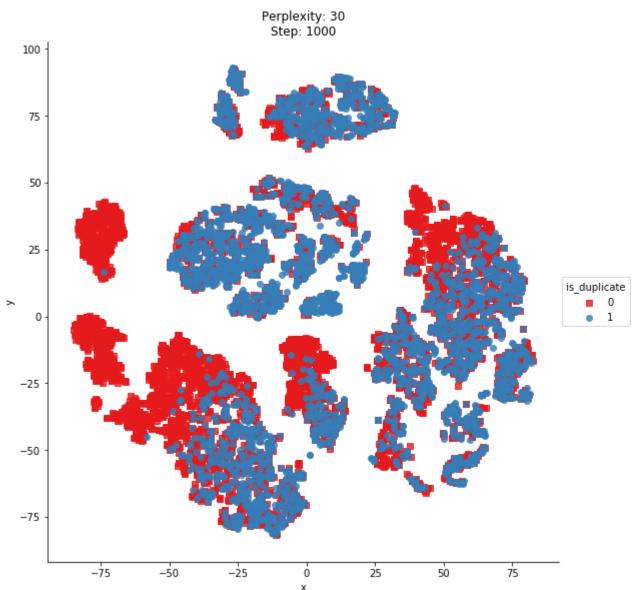
[t-SNE] Iteration 950: error = 1.2231647, gradient norm = 0.0000590 (50 iterations in 18.947s) [t-SNE] Iteration 1000: error = 1.2148699, gradient norm = 0.0000562 (50 iterations in 18.798s)

In [108]: tsne\_2D\_data.shape

[t-SNE] Error after 1000 iterations: 1.214870

Total time taken to get the tsne 2D data: 5.97 minutes

Out[108]: (10000, 2)



Analysis: The overlap is still existing. We will try and see if the overlap i resolved by 3D t-SNE Scatter Plot.

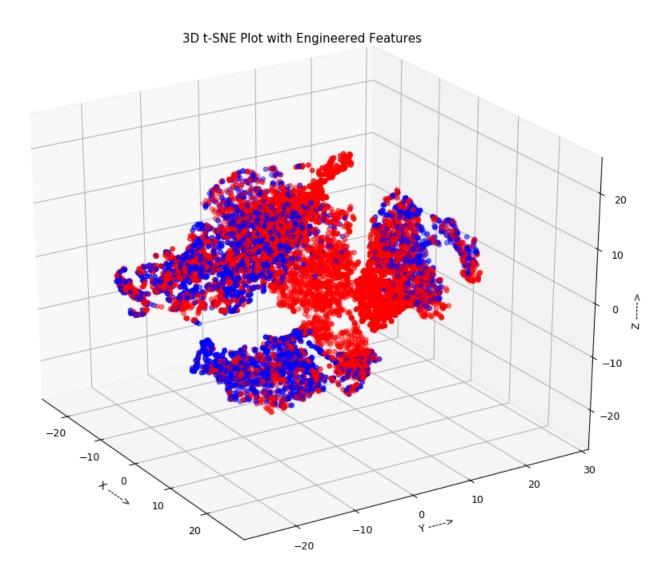
```
3D t-SNE Plot
 In [111]: # We will take n components=3, because we want to reduce the dimensions to 3D.
            tsne 3D model = TSNE(
                n components=3, init='random', method='barnes hut', n iter=1000,
                verbose=2, angle=0.5
            st = time.time()
            tsne 3D data = tsne 3D model.fit transform(df sample normalized)
            print('Total time taken to get the tsne 3D data: {} minutes'
                  .format(round((time.time()-st)/60, 2)))
            [t-SNE] Computing 91 nearest neighbors...
            [t-SNE] Indexed 10000 samples in 0.048s...
            [t-SNE] Computed neighbors for 10000 samples in 1.385s...
            [t-SNE] Computed conditional probabilities for sample 1000 / 10000
            [t-SNE] Computed conditional probabilities for sample 2000 / 10000
            [t-SNE] Computed conditional probabilities for sample 3000 / 10000
            [t-SNE] Computed conditional probabilities for sample 4000 / 10000
            [t-SNE] Computed conditional probabilities for sample 5000 / 10000
            [t-SNE] Computed conditional probabilities for sample 6000 / 10000
            [t-SNE] Computed conditional probabilities for sample 7000 / 10000
            [t-SNE] Computed conditional probabilities for sample 8000 / 10000
            [t-SNE] Computed conditional probabilities for sample 9000 / 10000
            [t-SNE] Computed conditional probabilities for sample 10000 / 10000
            [t-SNE] Mean sigma: 0.106319
            [t-SNE] Computed conditional probabilities in 0.606s
            [t-SNE] Iteration 50: error = 96.8366852, gradient norm = 0.0102541 (50 iterations in 46.923s)
            [t-SNE] Iteration 100: error = 76.2636795, gradient norm = 0.0021178 (50 iterations in 23.152s)
            [t-SNE] Iteration 150: error = 73.9572372, gradient norm = 0.0009611 (50 iterations in 20.420s)
            [t-SNE] Iteration 200: error = 73.1085739, gradient norm = 0.0006543 (50 iterations in 22.001s)
            [t-SNE] Iteration 250: error = 72.6302032, gradient norm = 0.0005056 (50 iterations in 23.378s)
            [t-SNE] KL divergence after 250 iterations with early exaggeration: 72.630203
            [t-SNE] Iteration 300: error = 2.2071168, gradient norm = 0.0007999 (50 iterations in 25.768s)
            [t-SNE] Iteration 350: error = 1.7142066, gradient norm = 0.0002601 (50 iterations in 29.803s)
            [t-SNE] Iteration 400: error = 1.4810127, gradient norm = 0.0001295 (50 iterations in 30.031s)
            [t-SNE] Iteration 450: error = 1.3485044, gradient norm = 0.0000788 (50 iterations in 29.573s)
            [t-SNE] Iteration 500: error = 1.2653854, gradient norm = 0.0000545 (50 iterations in 29.798s)
            [t-SNE] Iteration 550: error = 1.2097586, gradient norm = 0.0000430 (50 iterations in 28.944s)
            [t-SNE] Iteration 600: error = 1.1726710, gradient norm = 0.0000362 (50 iterations in 30.047s)
            [t-SNE] Iteration 650: error = 1.1479626, gradient norm = 0.0000337 (50 iterations in 29.632s)
```

[t-SNE] Error after 1000 iterations: 1.086666 Total time taken to get the tsne 3D data: 9.88 minutes

[t-SNE] Iteration 700: error = 1.1322137, gradient norm = 0.0000325 (50 iterations in 31.174s) [t-SNE] Iteration 750: error = 1.1215048, gradient norm = 0.0000307 (50 iterations in 31.177s) [t-SNE] Iteration 800: error = 1.1133183, gradient norm = 0.0000262 (50 iterations in 32.443s) [t-SNE] Iteration 850: error = 1.1059847, gradient norm = 0.0000245 (50 iterations in 31.988s) [t-SNE] Iteration 900: error = 1.0989758, gradient norm = 0.0000231 (50 iterations in 29.923s) [t-SNE] Iteration 950: error = 1.0922613, gradient norm = 0.0000214 (50 iterations in 32.160s) [t-SNE] Iteration 1000: error = 1.0866659, gradient norm = 0.0000201 (50 iterations in 32.592s) 10000

Out[118]: (10000,)

```
In [121]: %matplotlib notebook
          fig_tsne_3D = plt.figure(figsize=(10,8))
          ax_tsne = Axes3D(fig_tsne_3D)
          # The following lists 'color' and 'l' are already defined
          # above while plotting for 3D PCA as the following:
          # color = ['red', 'blue']
          # L = [float(i) for i in sample_labels]
          ax_tsne.scatter(
              tsne_3D_data[:,0:1], tsne_3D_data[:,1:2],
              tsne_3D_data[:,2:3], c = 1,
              cmap = matplotlib.colors.ListedColormap(color)
          ax_tsne.set_title('3D t-SNE Plot with Engineered Features')
          ax_tsne.set_xlabel('X ---->')
          ax_tsne.set_ylabel('Y ---->')
          ax_tsne.set_zlabel('Z ---->')
          plt.show()
```



Analysis: There's a clear overlap of points in all the red (is\_duplicate=0) & blue (is\_duplicate=1) points. In the middle, we can see that the red points are nicely seggregated in the 3D t-SNE plot. These advanced engineered features are definitely useful.

The use of the features will be further enhanced when we use tf-idf weighted word2vec to convert 'question1' and 'question2' into numerical word vectors, which retain the similarity depending on the sentiment of the questions.