

QQP3 - Featurizing text data with TF-IDF weighted Average Word2Vec

April 27, 2018

0.1 3.6 Featurizing text data with TF-IDF weighted word-vectors and Avg.word-vectors

In [1]: *# We will only operate on the first 100k points as 8GB RAM is not enough.*

```
# Library imports:
import numpy as np
import pandas as pd
from time import time
import warnings
warnings.filterwarnings('ignore')
import matplotlib.pyplot as plt
import sys
import os
from tqdm import tqdm

# We can extract word2vec vectors using spacy
# If there are any dependency issues, please folow these links:
# https://github.com/explosion/spaCy/issues/1721
# http://landinghub.visualstudio.com/visual-cpp-build-tools
import spacy
```

In [2]: *# We will use the following function to time our code:*

```
def time_taken(start_time):
    print("> Time taken:",
          round(time()-start_time, 2), "seconds")
    return
```

In []: *# We will import more libraries as and when required.*

In [5]: st = time()

```
# Import sample from the original dataset:
df = pd.read_csv("../train/train.csv", nrows=100000)

# Encode all the questions to unicode format:
df['question1'] = df['question1'].apply(lambda x: str(x))
```

```

df['question2'] = df['question2'].apply(lambda x: str(x))

time_taken(st)
df.shape

~> Time taken: 0.42 seconds

Out[5]: (100000, 6)

In [6]: df.head(2)

Out[6]:
```

	id	qid1	qid2	question1	\
0	0	1	2	What is the step by step guide to invest in sh...	
1	1	3	4	What is the story of Kohinoor (Koh-i-Noor) Dia...	

	question2	is_duplicate
0	What is the step by step guide to invest in sh...	0
1	What would happen if the Indian government sto...	0

0.1.1 3.6.1 Computing TF-IDF weighted Average Word2Vec Vectors

```

In [7]: from sklearn.feature_extraction.text import TfidfVectorizer

# Merge texts into a single list:
questions = list(df['question1']) + list(df['question2'])

# Create TfidfVectorizer instance:
tfidf = TfidfVectorizer(lowercase=False)

# Get the parameters in tfidf instance:
print(tfidf)

TfidfVectorizer(analyzer='word', binary=False, decode_error='strict',
dtype=<class 'numpy.int64'>, encoding='utf-8', input='content',
lowercase=False, max_df=1.0, max_features=None, min_df=1,
ngram_range=(1, 1), norm='l2', preprocessor=None, smooth_idf=True,
stop_words=None, strip_accents=None, sublinear_tf=False,
token_pattern='(?u)\\b\\w+\\b', tokenizer=None, use_idf=True,
vocabulary=None)

In [8]: st =time()

# Now apply tfidf transform:
tfidf.fit_transform(questions)

time_taken(st)

```

~> Time taken: 3.13 seconds

```
In [9]: st = time()
```

```
# We will now take all the tf-idf vectored values into a dictionary:  
# key:word and value:tf-idf score  
word2tfidf = dict(zip(tfidf.get_feature_names(), tfidf.idf_))  
  
time_taken(st)
```

~> Time taken: 0.08 seconds

- We have the TF-IDF Scores. We will now convert each question to a weighted average of word2vec vectors by using these TF-IDF Scores.
- We use a pre-trained GLoVe model which comes free with spacy library. The model itself is trained on Wikipedia data. More about the pre-trained model @ <https://spacy.io/usage/vectors-similarity>
- Because it is trained on Wikipedia, the model is strong in terms of word semantics

```
In [ ]: # We can either load and use 'en_vectors_web_lg' or 'en_core_web_sm'.  
# The difference between the two is that, 'en_core_web_sm' is a smaller  
# model compared to 'en_vectors_web_lg'  
  
# If the system has less RAM, it is better to load 'en_core_web_sm' model  
# as the word-vector. The only disadvantage with a smaller model is that  
# it will give us similarity vectors with lesser accuracy.  
  
# Note: there are 2 full version models of word2vec models trained on  
# wikipedia data. One of them is 'en_core_web_lg' and the other one is  
# 'en_vectors_web_lg'. If we want to make changes to the word2vec model,  
# then we can load the 'en_vectors_web_lg' word2vec model, otherwise, we  
# go with 'en_core_web_lg' word2vec model.  
  
# Note: We can't load these pre-trained word2vec models unless and until  
# we download the models through the console of the system.  
  
#####  
# Syntax: python -m spacy download model_name #  
#####  
  
# Example: python -m spacy download en_core_web_lg
```

```
In [10]: # Loading the smaller model:  
nlp = spacy.load('en_core_web_sm')  
st = time()  
  
vector1 = []
```

```

# https://github.com/noamraph/tqdm
# tqdm is used to print the progress bar.
for q1 in tqdm(list(df.question1)):
    doc = nlp(q1)
    # 384 dimensional vector
    mean_vec1 = np.zeros([len(doc), 384])
    for word in doc:
        # word2vec:
        vec1 = word.vector

        # Fetch the tf-idf score:
        try:
            idf = word2tfidf[str(word)]
        except:
            idf = 0

        # Compute final vector:
        mean_vec1 += vec1 * idf
    mean_vec1 = mean_vec1.mean(axis=0)
    vector1.append(mean_vec1)
df['q1_feats_m'] = list(vector1)

```

100%|| 100000/100000 [19:04<00:00, 87.35it/s]

In [11]: time_taken(st)

~> Time taken: 1144.88 seconds

In [12]: st = time()

```

vector2 = []
for q2 in tqdm(list(df.question2)):
    doc = nlp(q2)
    mean_vec2 = np.zeros([len(doc), 384])
    for word in doc:
        # word2vec:
        vec2 = word.vector
        # Fetch idf score:
        try:
            idf = word2tfidf[str(word)]
        except:
            idf = 0

        # Compute final vector:
        mean_vec2 += vec2 * idf
    mean_vec2 = mean_vec2.mean(axis=0)
    vector2.append(mean_vec2)

```

```
df['q2_feats_m'] = list(vector2)

time_taken(st)

100%|| 100000/100000 [19:29<00:00, 85.47it/s]
```

~> Time taken: 1170.01 seconds

In [13]: *# Get the files in the current working directory*

```
files_in_cwd = os.listdir()
index = 0
print("<idx>. <Filename>")
for f in files_in_cwd:
    print("{} . {}".format(index, f))
    index += 1
```

<idx>. <Filename>

```
0. .ipynb_checkpoints
1. df_fe_without_preprocessing_train.csv
2. nlp_features_train.csv
3. QQP1 - BasicEDA, TextPreprocessing, BasicFeaturization, AdvancedFeaturization (NLP & Fuzzy :
4. QQP1.py
5. QQP2 - Word Cloud, PCA & t-SNE 2D and 3D Visualizations of Engineered Features.ipynb
6. QQP2.py
7. QQP3 - Featurizing text data with TF-IDF weighted Average Word2Vec.ipynb
8. quora.png
9. train_dup_question_pairs.txt
10. train_non_dup_question_pairs.txt
```

In [14]: *# We will read some previously saved .csv files like:*

```
# 1. nlp_features_train.csv
# 2. df_fe_without_preprocessing_train.csv
# and we will use the data generated now and finally merge all of the data into
# a single pandas dataframe. The dataframe size may get really large.
```

```
st = time()
# Load the nlp_features_train.csv file into a dataframe:
if os.path.isfile('nlp_features_train.csv'):
    df_nlp = pd.read_csv("nlp_features_train.csv", encoding='latin-1', nrows=100000)
else:
    print('Generate the file by running the code in QQP1.')

# Load the df_fe_without_preprocessing_train.csv file into a dataframe:
if os.path.isfile('df_fe_without_preprocessing_train.csv'):
    df_pre = pd.read_csv('df_fe_without_preprocessing_train.csv', encoding=\
        'latin-1', nrows=100000)
```

```

else:
    print('Generate the file by running the code in QQP1.')

time_taken(st)

~> Time taken: 1.55 seconds

In [15]: print(df_nlp.shape)
print(df_pre.shape)
# We will drop the unnecessary features and only keep the required ones:
df1 = df_nlp.drop(['qid1', 'qid2', 'question1', 'question2'], axis=1)

# df1 corresponds to advanced nlp and fuzzy engineered features:
df1.head()

(100000, 21)
(100000, 17)

Out[15]:
   id  is_duplicate  cwc_min  cwc_max  csc_min  csc_max  ctc_min  \
0    0             0  0.999980  0.833319  0.999983  0.999983  0.916659
1    1             0  0.799984  0.399996  0.749981  0.599988  0.699993
2    2             0  0.399992  0.333328  0.399992  0.249997  0.399996
3    3             0  0.000000  0.000000  0.000000  0.000000  0.000000
4    4             0  0.399992  0.199998  0.999950  0.666644  0.571420

      ctc_max  last_word_eq  first_word_eq  abs_len_diff  mean_len  \
0  0.785709             0.0             1.0             0.0      13.0
1  0.466664             0.0             1.0             0.0      12.5
2  0.285712             0.0             1.0             0.0      12.0
3  0.000000             0.0             0.0             0.0      12.0
4  0.307690             0.0             1.0             0.0      10.0

      token_set_ratio  token_sort_ratio  fuzz_ratio  fuzz_parital_ratio  \
0                100                93            93                100
1                 86                63            66                 75
2                 63                63            43                 47
3                 28                24             9                 14
4                 67                47            35                 56

      longest_substr_ratio
0                0.982759
1                0.596154
2                0.166667
3                0.039216
4                0.175000

In [16]: df2 = df_pre.drop(['qid1', 'qid2', 'question1', 'question2', 'is_duplicate'],\
                             axis=1)

```

```
# df2 corresponds to basic engineered features:
df2.head()
```

```
Out[16]:
```

	id	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words	\
0	0	1	1	66	57	14	12	
1	1	4	1	51	88	8	13	
2	2	1	1	73	59	14	10	
3	3	1	1	50	65	11	9	
4	4	3	1	76	39	13	7	

	word_Common	word_Total	word_share	freq_q1+q2	freq_q1-q2
0	10.0	23.0	0.434783	2	0
1	4.0	20.0	0.200000	5	3
2	4.0	24.0	0.166667	2	0
3	0.0	19.0	0.000000	2	0
4	2.0	20.0	0.100000	4	2

```
In [17]: # our original dataset with some additional features:
df.head()
```

```
Out[17]:
```

	id	qid1	qid2	question1	\
0	0	1	2	What is the step by step guide to invest in sh...	
1	1	3	4	What is the story of Kohinoor (Koh-i-Noor) Dia...	
2	2	5	6	How can I increase the speed of my internet co...	
3	3	7	8	Why am I mentally very lonely? How can I solve...	
4	4	9	10	Which one dissolve in water quikly sugar, salt...	

	question2	is_duplicate	\
0	What is the step by step guide to invest in sh...	0	
1	What would happen if the Indian government sto...	0	
2	How can Internet speed be increased by hacking...	0	
3	Find the remainder when 23^{24} i...	0	
4	Which fish would survive in salt water?	0	

	q1_feats_m	\
0	[122.490798712, 100.359120488, 72.0331508666, ...	
1	[-74.5846772194, 53.8620963991, 81.0885115862,...	
2	[-5.10626339912, 73.7096084356, 14.3268437684,...	
3	[5.90131050348, -34.4693912566, 48.9884575009,...	
4	[48.4207775295, 38.2941785157, 121.9611063, 54...	

	q2_feats_m
0	[126.564217329, 96.0618406534, 42.2021160275, ...
1	[-105.099983424, 79.1588504314, 77.5340879094,...
2	[6.49532223493, 16.2452982366, 2.65493392944, ...
3	[38.9078674316, 43.9539289773, -24.3469197154,...
4	[31.6172962189, 62.5719087124, 1.96994256973, ...

```
In [18]: # We will drop ['qid1', 'qid2', 'question1', 'question2', 'is_duplicate'] from df:
df3 = df.drop(['qid1', 'qid2', 'question1', 'question2', 'is_duplicate'], axis=1)
df3.head()
```

```
Out [18]:
```

	id	q1_feats_m \
0	0	[122.490798712, 100.359120488, 72.0331508666, ...
1	1	[-74.5846772194, 53.8620963991, 81.0885115862, ...
2	2	[-5.10626339912, 73.7096084356, 14.3268437684, ...
3	3	[5.90131050348, -34.4693912566, 48.9884575009, ...
4	4	[48.4207775295, 38.2941785157, 121.9611063, 54...

	q2_feats_m
0	[126.564217329, 96.0618406534, 42.2021160275, ...
1	[-105.099983424, 79.1588504314, 77.5340879094, ...
2	[6.49532223493, 16.2452982366, 2.65493392944, ...
3	[38.9078674316, 43.9539289773, -24.3469197154, ...
4	[31.6172962189, 62.5719087124, 1.96994256973, ...

```
In [20]: # q1_feats_m has each row as a list. Therefore, we will extract it into a
# dataframe as:
st = time()
df_q1 = pd.DataFrame(df3.q1_feats_m.values.tolist(), index = df3.index)
time_taken(st)
df_q1.head()
```

~> Time taken: 14.04 seconds

```
Out [20]:
```

	0	1	2	3	4	5	\
0	122.490799	100.359120	72.033151	115.891096	-48.144981	34.736722	
1	-74.584677	53.862096	81.088512	98.550397	-50.356915	53.286622	
2	-5.106263	73.709608	14.326844	104.493053	1.258413	35.409146	
3	5.901311	-34.469391	48.988458	59.481399	40.695803	-41.397960	
4	48.420778	38.294179	121.961106	54.678226	-45.466374	38.553049	

	6	7	8	9	...	374	\
0	-172.386330	-93.059744	113.417203	51.259765	...	12.462868	
1	-37.665547	-82.297257	45.744834	-8.385913	...	-21.548015	
2	-149.265339	-97.636930	42.259155	51.435161	...	3.012211	
3	-36.726121	24.031034	0.295455	-29.501785	...	13.059348	
4	-294.462586	-105.776589	103.886341	65.766421	...	13.320748	

	375	376	377	378	379	380	381	\
0	41.063396	8.037371	-15.198150	18.056487	6.217941	-30.221076	3.659344	
1	-11.906959	20.344241	1.829228	-16.460159	-5.656435	-10.035233	-4.768943	
2	14.140741	-2.977540	-3.214739	4.373585	2.911802	-20.323167	9.798284	
3	1.411459	-1.874297	-7.867466	17.947856	12.057635	-10.482685	5.230752	
4	42.630676	11.245030	-21.892262	43.775802	8.189654	-34.812249	8.047953	

	382	383
0	-1.687294	-1.825006
1	-12.692666	-5.208524
2	11.907082	-8.814535
3	10.150245	5.845988
4	9.497889	5.378521

[5 rows x 384 columns]

In [21]: *# q12_feats_m has each row as a list. Therefore, we will extract it into a # dataframe as:*

```
st = time()
df_q2 = pd.DataFrame(df3.q2_feats_m.values.tolist(), index = df3.index)
time_taken(st)
df_q2.head()
```

~> Time taken: 14.85 seconds

Out [21]:

	0	1	2	3	4	5	\
0	126.564217	96.061841	42.202116	95.969994	-37.314082	39.737327	
1	-105.099983	79.158850	77.534088	58.330385	-41.438078	115.591662	
2	6.495322	16.245298	2.654934	86.827784	-34.626589	95.729673	
3	38.907867	43.953929	-24.346920	86.120009	0.079079	-9.801455	
4	31.617296	62.571909	1.969943	36.472732	-45.163165	66.659808	

	6	7	8	9	...	374	\
0	-148.516119	-88.340872	110.552041	62.843040	...	16.188503	
1	-142.872375	-125.501038	23.816001	25.313954	...	-4.432317	
2	-123.613627	-115.022091	53.958783	61.496209	...	8.264448	
3	-60.949873	-37.361491	49.504973	-22.386544	...	3.488654	
4	-105.894651	-22.777562	59.957627	62.017545	...	-2.440844	

	375	376	377	378	379	380	\
0	33.233713	6.971700	-14.820828	15.534945	8.205955	-25.256606	
1	-4.367793	41.101273	-0.930737	-15.686246	-7.275999	2.756560	
2	-2.244750	11.084606	-16.741266	14.854023	15.726977	-1.298039	
3	3.906499	13.387563	-6.640244	6.378005	6.028185	2.511873	
4	11.887040	8.019029	-15.028031	8.280575	1.703147	-6.503707	

	381	382	383
0	1.552828	1.651827	0.267462
1	-7.351970	3.103773	0.440425
2	14.340431	11.669012	10.423255
3	-3.830347	5.421078	6.161891
4	11.263387	11.556818	2.500520

[5 rows x 384 columns]

```
In [23]: print("Number of features in nlp dataframe:", df1.shape[1])
        print("Number of features in preprocessed dataframe:", df2.shape[1])
        print("Number of features in question1 w2v dataframe:", df_q1.shape[1])
        print("Number of features in question2 w2v dataframe:", df_q2.shape[1])
        print("Number of features in the final dataframe:",\
              df1.shape[1] + df2.shape[1] + df_q1.shape[1] + df_q2.shape[1])
```

```
Number of features in nlp dataframe: 17
Number of features in preprocessed dataframe: 12
Number of features in question1 w2v dataframe: 384
Number of features in question2 w2v dataframe: 384
Number of features in the final dataframe: 797
```

```
In [25]: st = time()
```

```
# The following code might take some time to execute, depending on the system
# configuration.
if not os.path.isfile('final_features_100k.csv'):

    # Attach 'id' attribute to question1 and question2 w2v vectors:
    df_q1['id'] = df1['id']
    df_q2['id'] = df1['id']

    # Merge nlp_features with preprocessing_features:
    df1 = df1.merge(df2, on='id', how='left')

    # Merge question1 and question2 w2v vectors and save them in df2 variable:
    df2 = df_q1.merge(df_q2, on='id', how='left')

    # We will now merge df1 and df2 into result:
    result = df1.merge(df2, on='id', how='left')

    # Save as a .csv file to use when applying k-NN to classify the points:
    result.to_csv('final_features_100k.csv')

time_taken(st)
```

```
~> Time taken: 0.0 seconds
```

final_features_100k.csv file is generated, which is of 1.37GB. The file will be used to apply k-NN, so that we will be able to know how accurate the k-NN classifier is.