

Quora Question Pairs

Sukanya Behera [160714733008]

Chandrabhatta Sriram [160713733042]

Abul Faiz Mohammed Abdul Hadi Mouzzam [160714733108]

Project Guide: E. Shailaja [Asst. Prof. Dept. of CSE]

Project Co-ordinator: R. Sandeep [Asst. Prof. Dept. of CSE]

Introduction [Problem Definition]

- More formally, the duplicate question detection problem can be defined as follows: given a pair of questions $q1$ and $q2$, train a model that learns the function:
 - $f(q1, q2) \rightarrow 0 \text{ or } 1$.
 - where 1 represents that $q1$ and $q2$ have the same intent and 0 otherwise.
 - NOTE: $q1$ and $q2$ are given as string data.

Problem Input [test.csv]

id	qid1	qid2	question1	question2
447	895	896	What are natural numbers?	What is a least natural number?
1518	3037	3038	Which pizzas are the most popularly ordered pizzas on Domino's menu?	How many calories does a Dominos pizza have?
3272	6542	6543	How do you start a bakery?	How can one start a bakery business?
3362	6722	6723	Should I learn python or Java first?	If I had to choose between learning Java and Python, what should I choose to learn first?

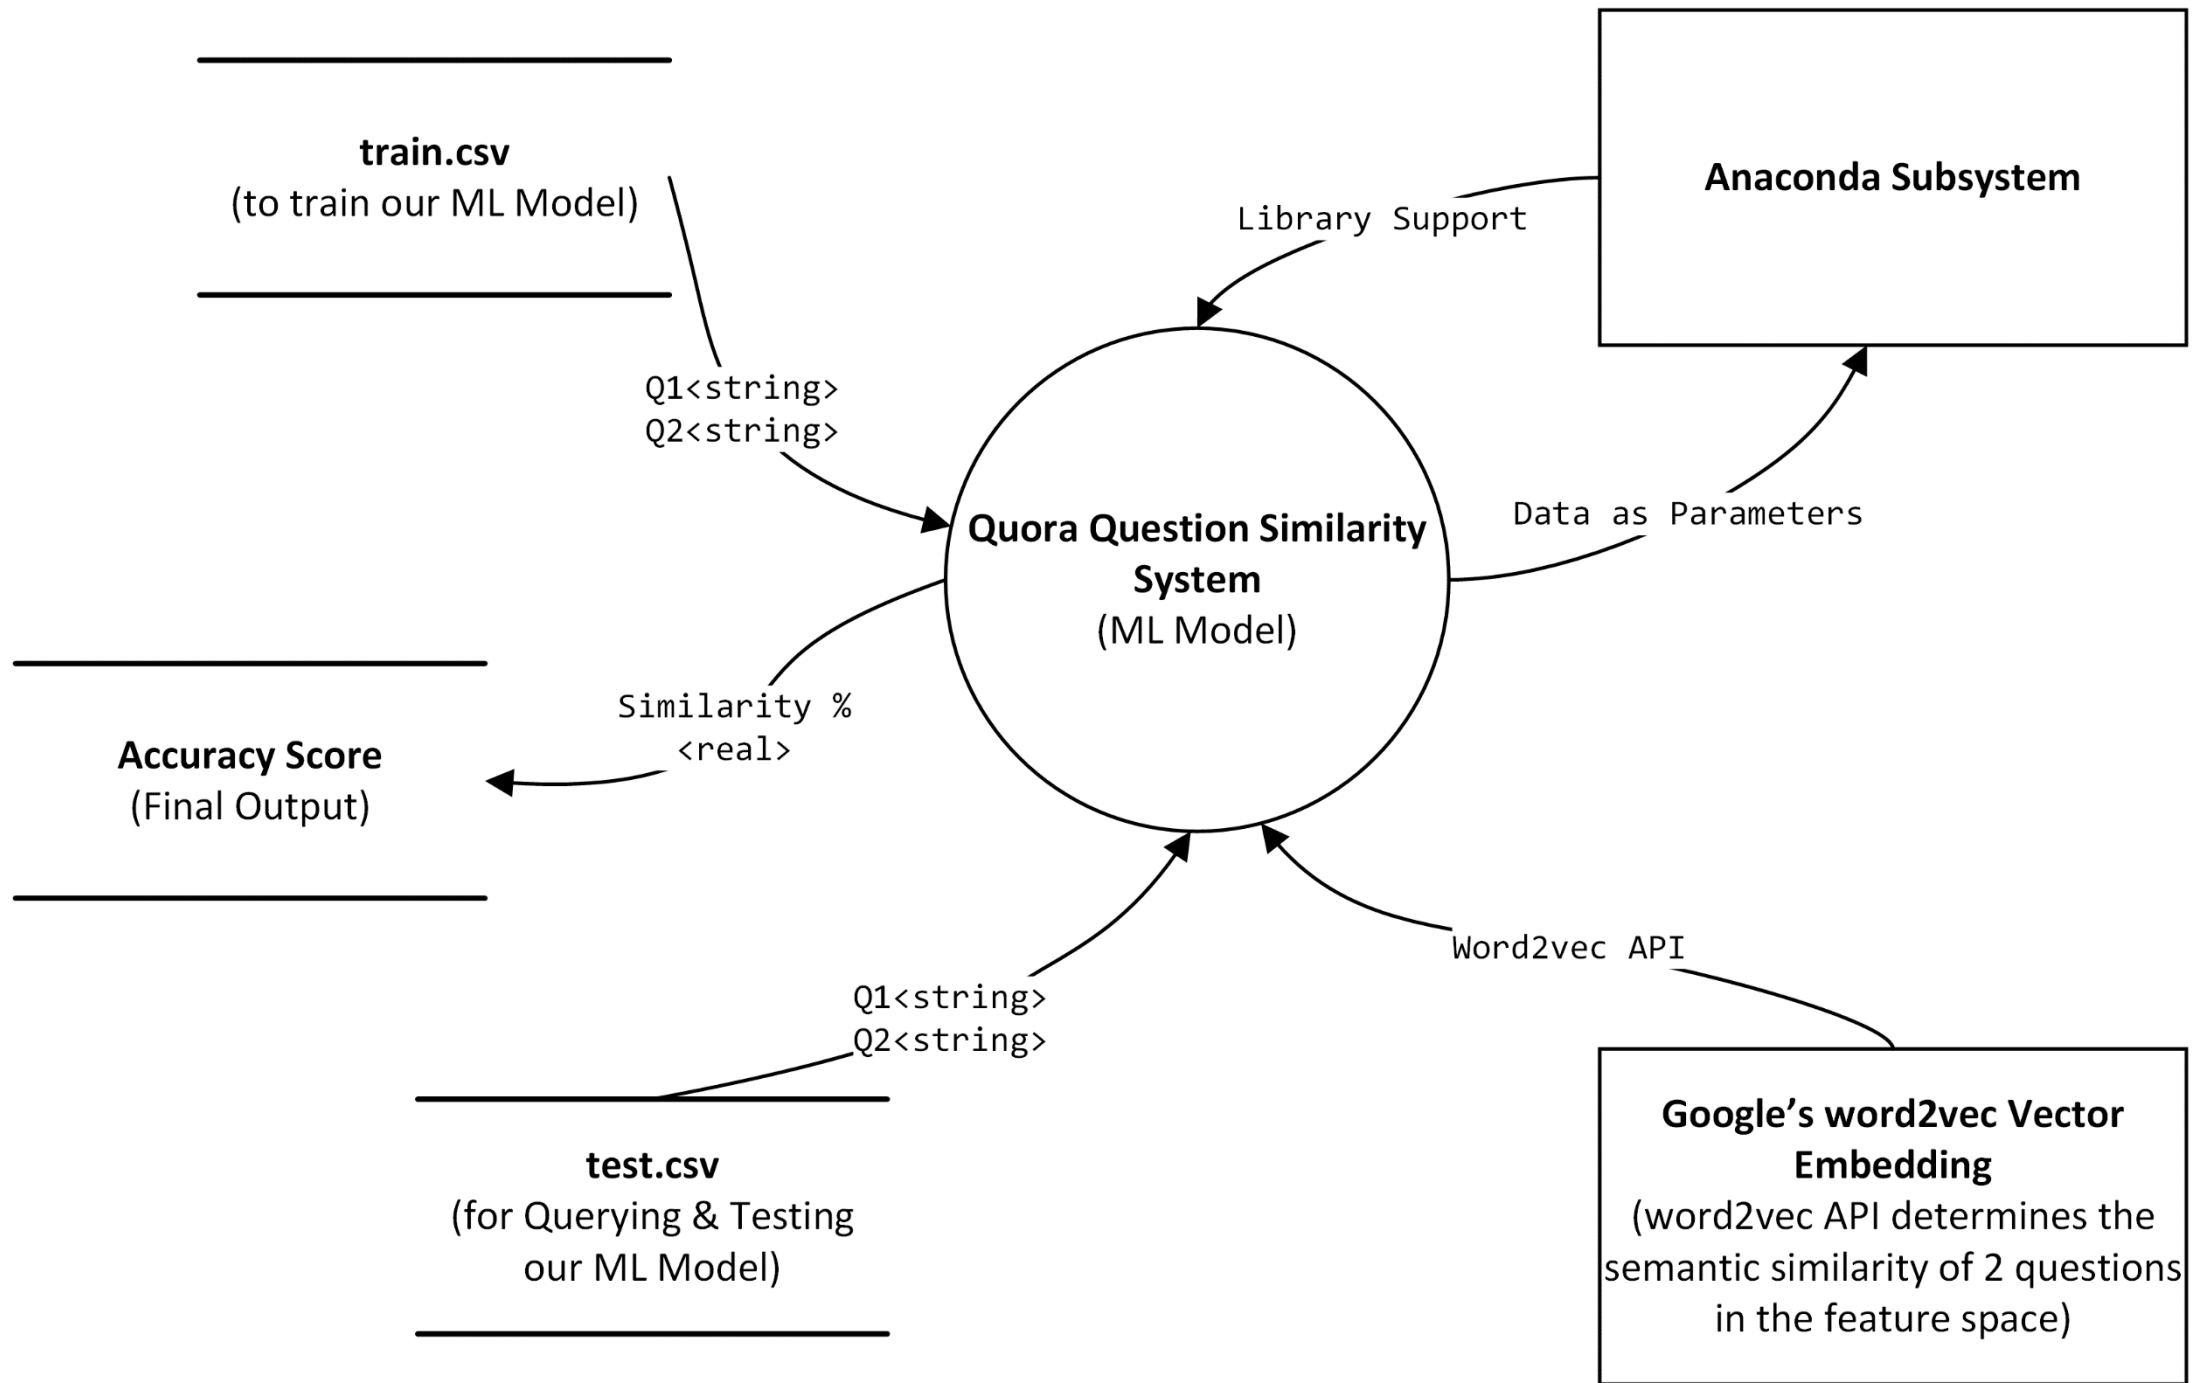
Problem Output [Expected]

id	qid1	qid2	question1	question2	is_duplicate
447	895	896	What are natural numbers?	What is a least natural number?	0
1518	3037	3038	Which pizzas are the most popularly ordered pizzas on Domino's menu?	How many calories does a Dominos pizza have?	0
3272	6542	6543	How do you start a bakery?	How can one start a bakery business?	1
3362	6722	6723	Should I learn python or Java first?	If I had to choose between learning Java and Python, what should I choose to learn first?	1

Modules

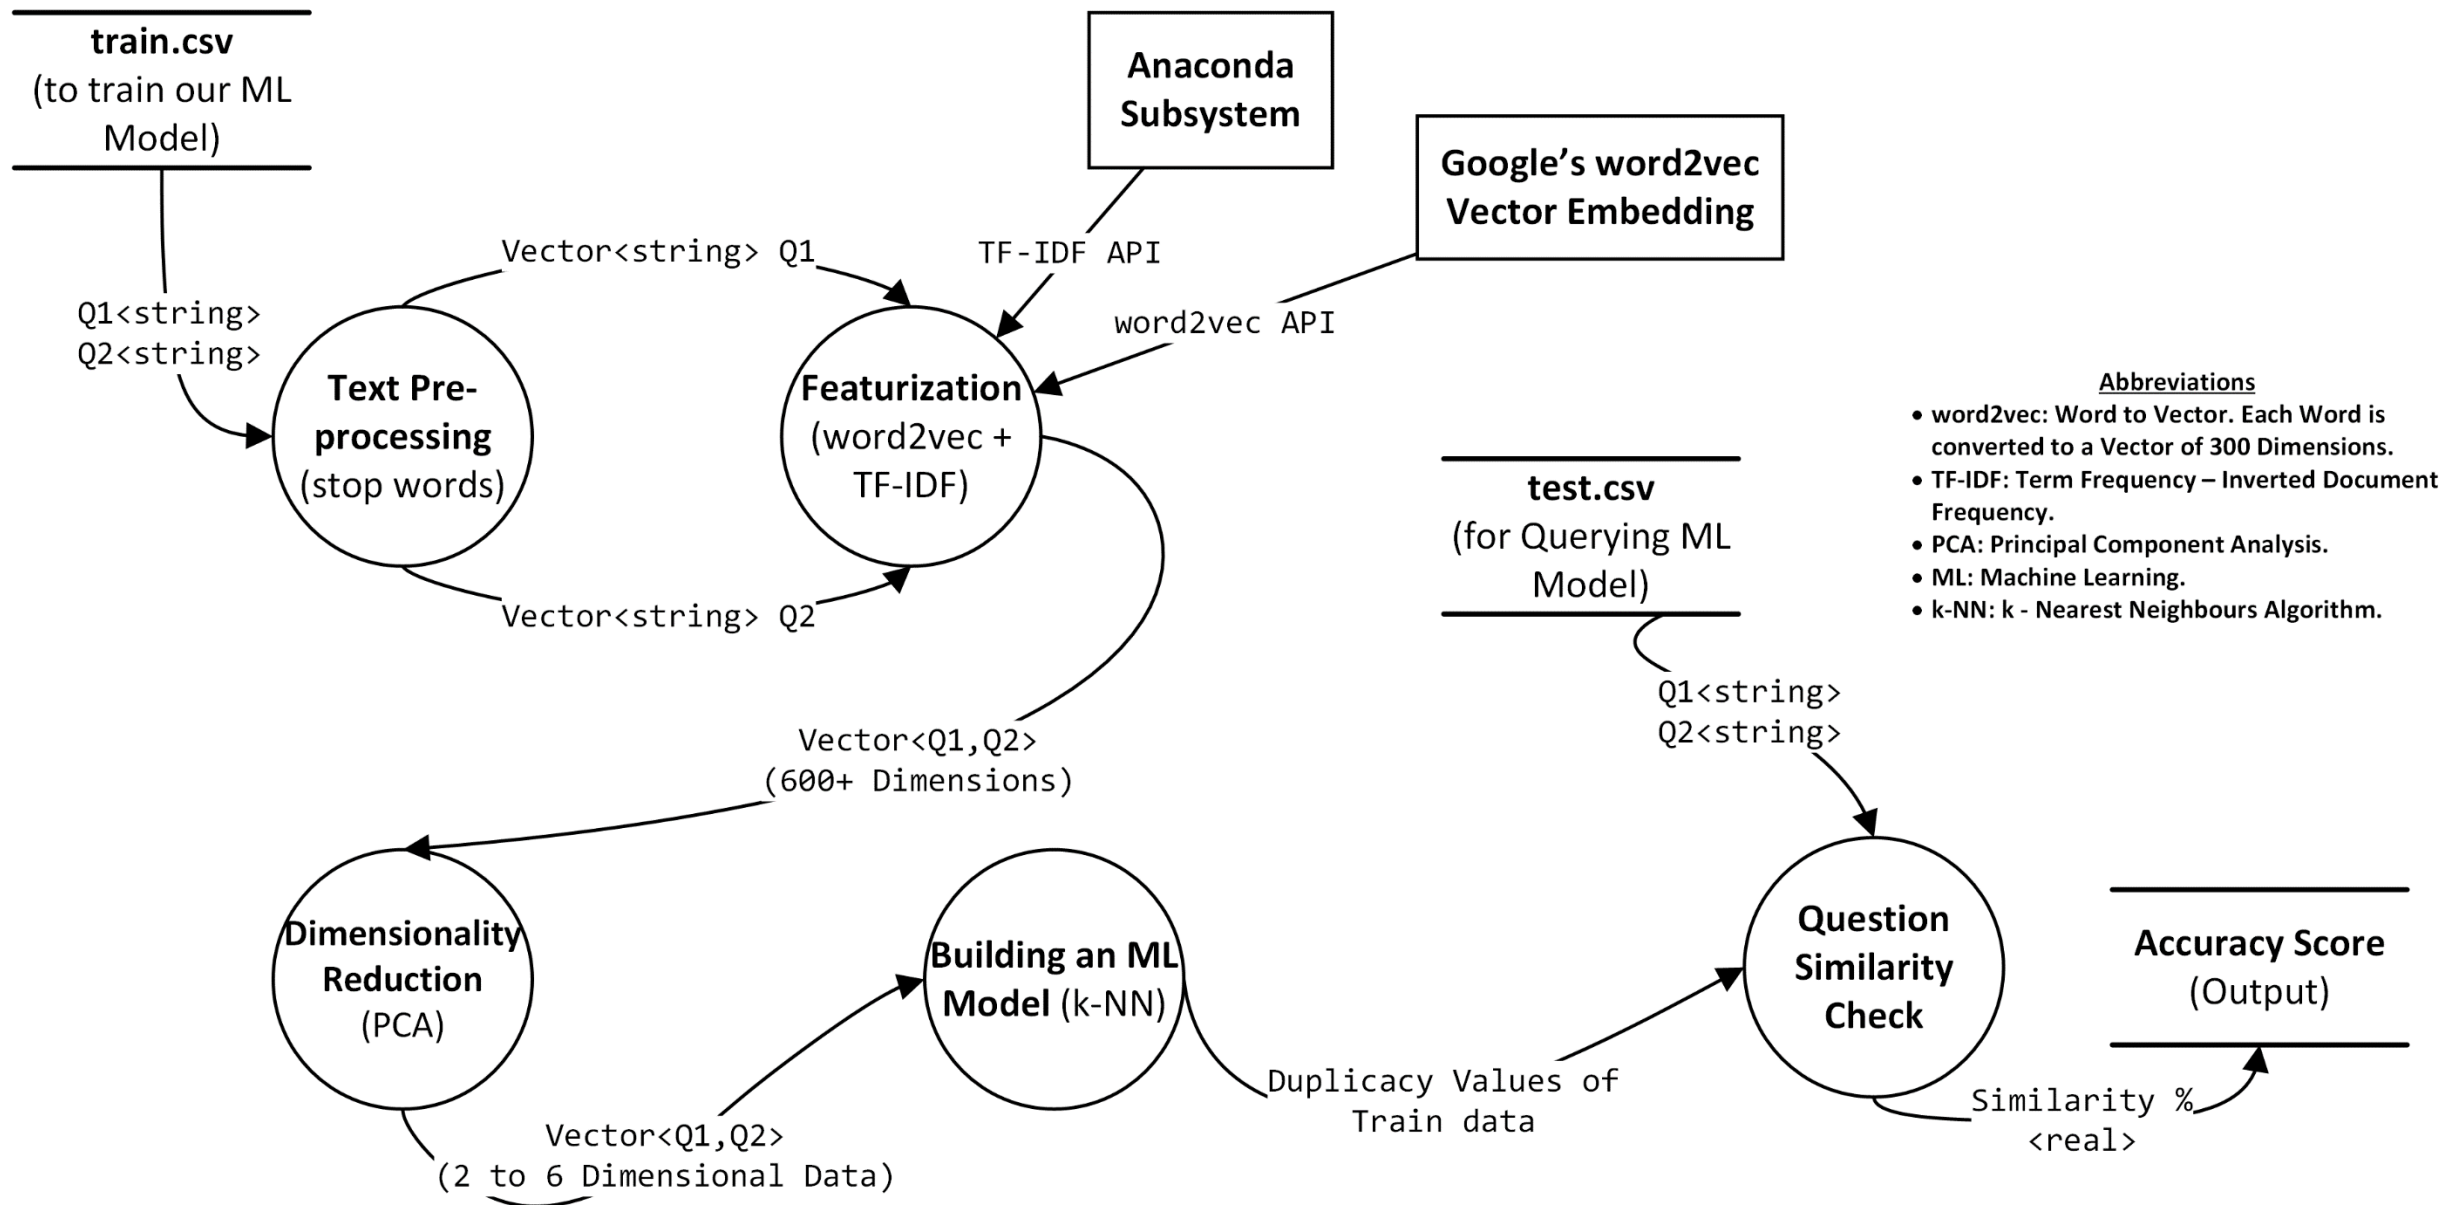
- Text Pre-processing: *text tokenization, uniform text casing, stop word removal, lemmatization, stemming*, etc.
- Featurization / Vectorization of the texts [$q1$ and $q2$]: Convert the given $q1$ and $q2$ into Word Vectors using word2vec. Word vectors are positioned in the vector space such that words that share common contexts in the corpus are located in close proximity to one another in the space.
- Apply a Machine Learning Model [using k – Nearest Neighbours algorithm] to get the Expected Output.

Data Flow Diagram – Level 0



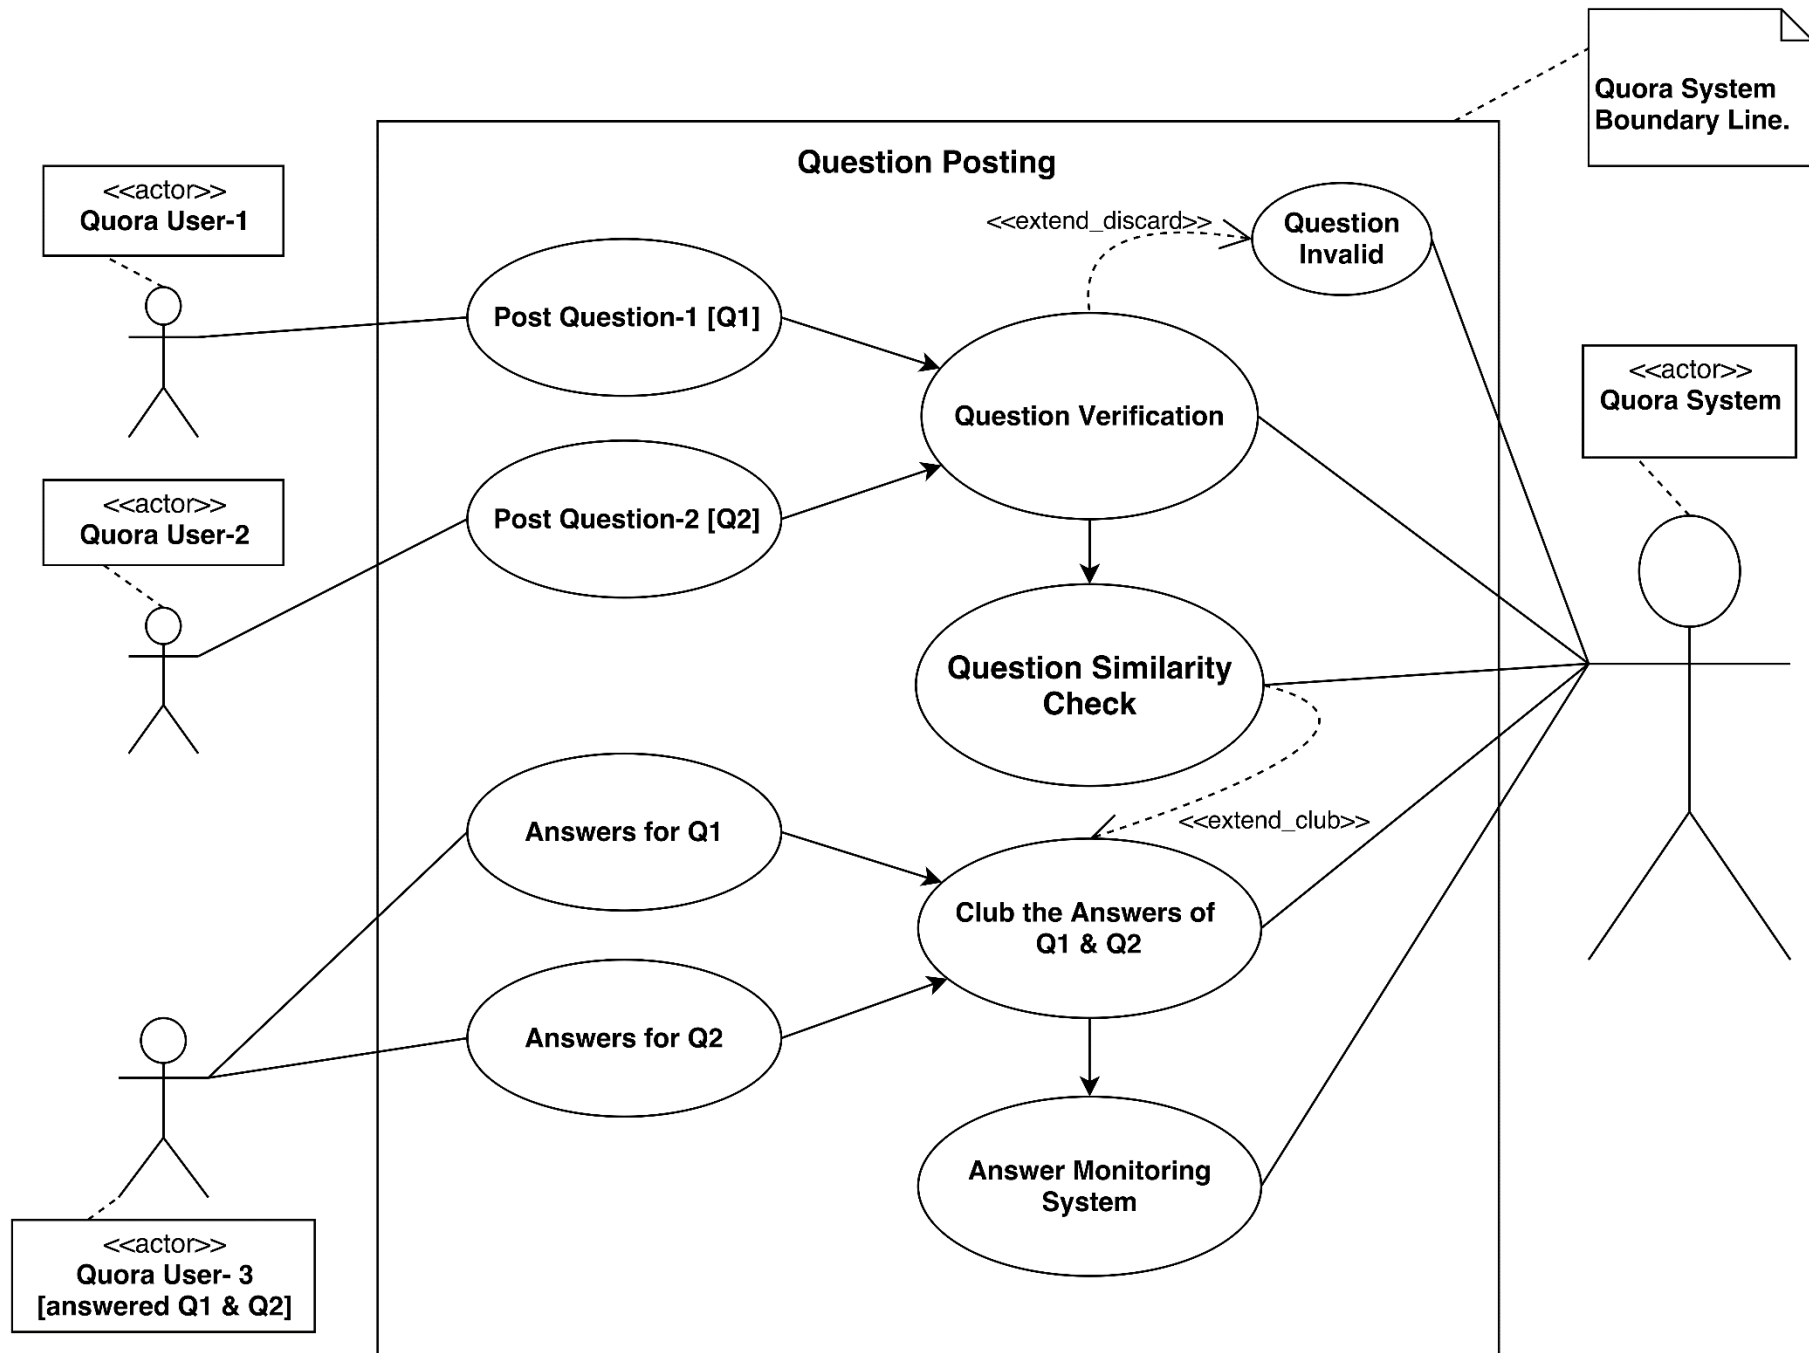
Data Flow Diagram – Level 1

Quora Question Similarity System



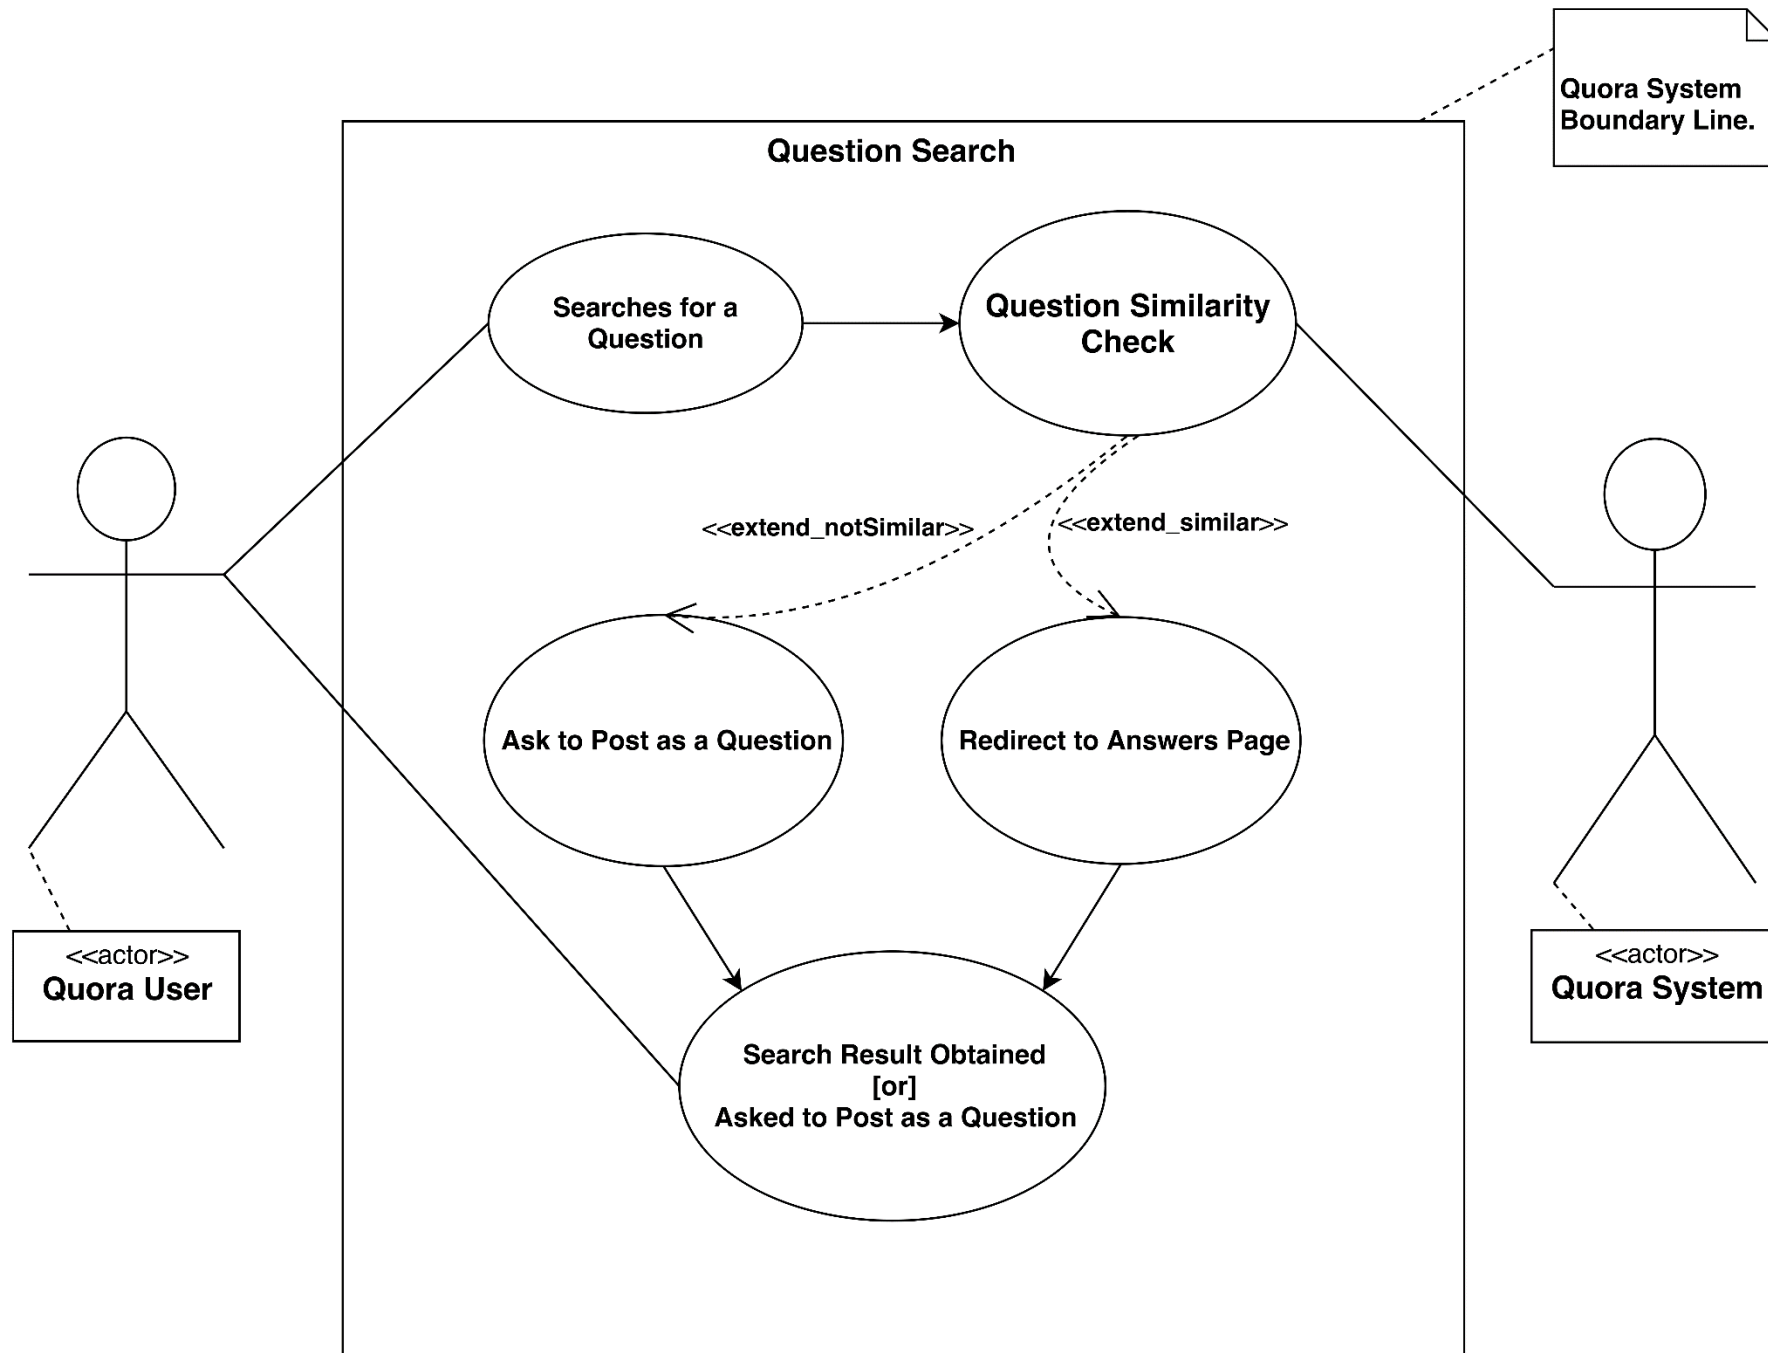
Use Case Diagram – 1

[Question Posting]




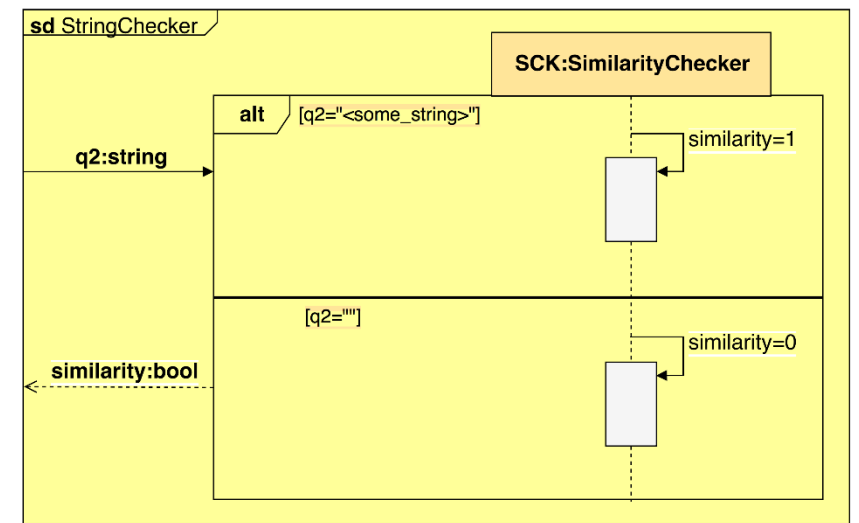
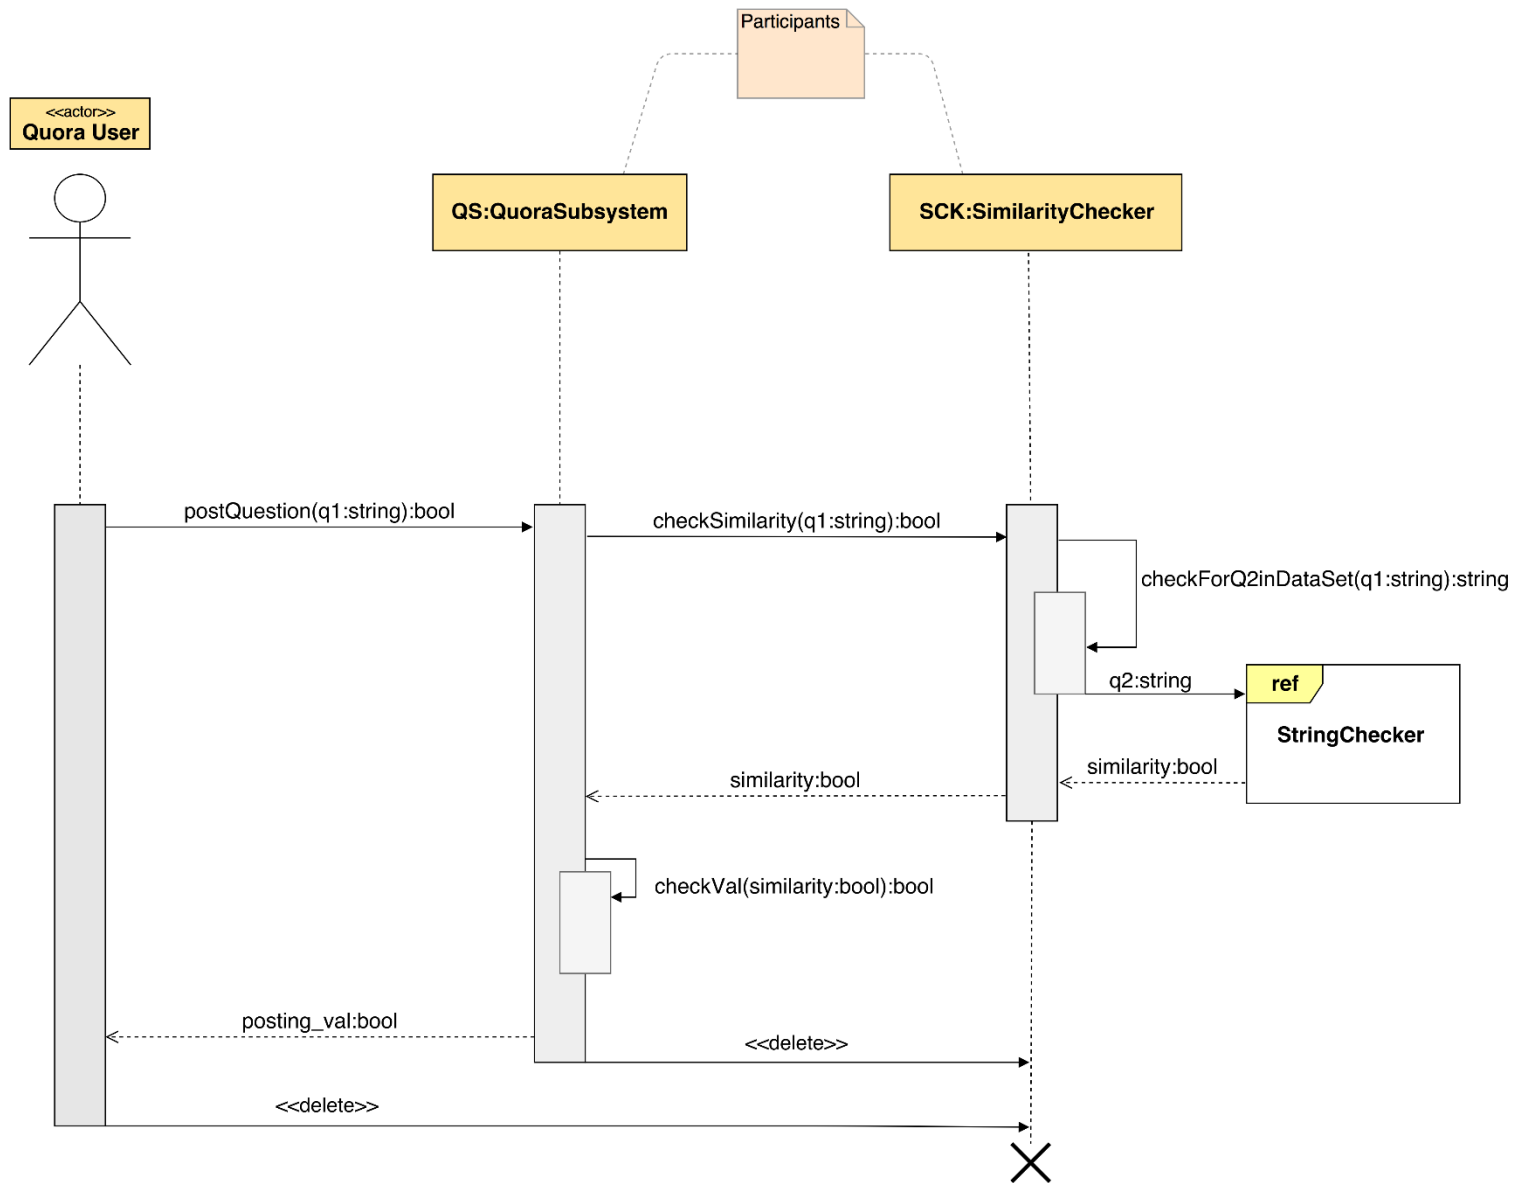
Use Case Diagram – 2

[Question Search]



Sequence Diagram






Deployment Diagram

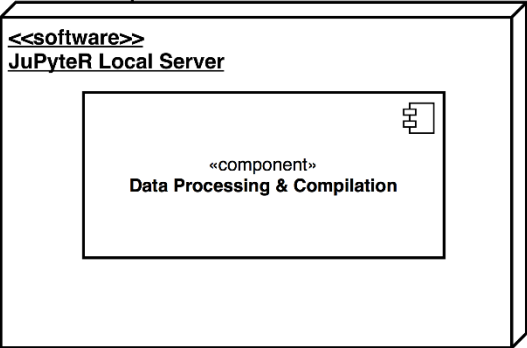
<<root_server>>
Quora_System

<<leaf_server>>
Question_Similarity_Subsystem




<<Data_Load>>

<<localhost_connection>>




<<Library_Support>>



<<HTTP>>

<<device>>
Client_Mobile/PC



Algorithm Information

- In ***k*-NN** classification, the output is a class membership. An object is classified by a majority vote of its neighbours, with the object being assigned to the class most common among its ***k*** nearest neighbours (***k*** is a positive integer, typically small). If ***k*** = 1, then the object is simply assigned to the class of that single nearest neighbour.
- Distance between 2 points in the feature space, can be found by taking Euclidean Distance (L^2 Norm) between those 2 points.

k-NN Brute Force Pseudo Code

Algorithm 1: Brute force k NN Algorithm

Input : \mathcal{Q} , a set query points and \mathcal{R} , a set of reference point;


Output: A list of k *reference* points for each *query* point;

```
1 foreach query point  $q \in \mathcal{Q}$  do
2   compute distances between  $q$  and all  $r \in \mathcal{R}$ ;
3   sort the computed distances;
4   select  $k$ -nearest reference points corresponding to  $k$  smallest distances;
```

Complexity: $O(mnd)$, which makes it a Polynomial Time Algorithm.


kNN Algorithm

0. Look at the data



Say you want to classify the grey point into a class. Here, there are three potential classes - lime green, green and orange.

1. Calculate distances



Start by calculating the distances between the grey point and all other points.

2. Find neighbours

Point Distance		
...	2.1	→ 1st NN
...	2.4	→ 2nd NN
...	3.1	→ 3rd NN
...	4.5	→ 4th NN

Next, find the nearest neighbours by ranking points by increasing distance. The nearest neighbours (NNs) of the grey point are the ones closest in dataspace.

3. Vote on labels

Class	# of votes	
	2	→ Class wins the vote! Point is therefore predicted to be of class .
	1	
	1	

Vote on the predicted class labels based on the classes of the k nearest neighbours. Here, the labels were predicted based on the $k=3$ nearest neighbours.


KD-Tree Algorithm for k-NN

```
function kdtree (list of points pointList, int depth)
{
    // Select axis based on depth so that axis cycles through all valid values
    var int axis := depth mod k;


    // Sort point list and choose median as pivot element
    select median by axis from pointList;

    // Create node and construct subtree
    node.location := median;
    node.leftChild := kdtree(points in pointList before median, depth+1);
    node.rightChild := kdtree(points in pointList after median, depth+1);
    return node;
}
```

KD-Tree Visualization for k-NN



***k*-d tree decomposition for the point set (2,3), (5,4), (9,6), (4,7), (8,1), (7,2).**



KD-Tree Algorithm Details

<i>k</i>-d tree		
Type	Multidimensional BST	
Invented	1975	
Invented by	Jon Louis Bentley	
Time complexity in big O notation		
Algorithm	Average	Worst case
Space	$O(n)$	$O(n)$
Search	$O(\log n)$	$O(n)$
Insert	$O(\log n)$	$O(n)$
Delete	$O(\log n)$	$O(n)$

Software & Hardware Requirements

- Software Requirements
 - OS: Windows 7/8/10, Linux (Cent/Ubuntu).
 - Packages: Python 3.6, Anaconda Package for Data Science & Machine Learning.
- Hardware Requirements:
 - Minimum System Requirements:
 - 8GB RAM, 5GB Disk Space, 1GHz Processor.
 - Recommended System Requirements:
 - 16GB RAM, 15GB Disk Space, 3.5GHz Multi Core Processor.

Code Snippets

- **freq_qid1** = Frequency of qid1's
- **freq_qid2** = Frequency of qid2's
- **q1len** = Length of q1
- **q2len** = Length of q2
- **q1_n_words** = Number of words in Question 1
- **q2_n_words** = Number of words in Question 2
- **word_Common** = (Number of common unique words in Question 1 and Question 2)
- **word_Total** = (Total num of words in Question 1 + Total num of words in Question 2)
- **word_share** = (word_common)/(word_Total)
- **freq_q1+q2** = sum total of frequency of qid1 and qid2
- **freq_q1-q2** = absolute difference of frequency of qid1 and qid2

```
In [14]: if os.path.isfile('df_fe_without_preprocessing_train.csv'):
        df = pd.read_csv("df_fe_without_preprocessing_train.csv",encoding='latin-1')
    else:
        df['freq_qid1'] = df.groupby('qid1')['qid1'].transform('count')
        df['freq_qid2'] = df.groupby('qid2')['qid2'].transform('count')
        df['q1len'] = df['question1'].str.len()
        df['q2len'] = df['question2'].str.len()
        df['q1_n_words'] = df['question1'].apply(lambda row: len(row.split(" ")))
        df['q2_n_words'] = df['question2'].apply(lambda row: len(row.split(" ")))

        def normalized_word_Common(row):
            w1 = set(map(lambda word: word.lower().strip(), row['question1'].split(" ")))
            w2 = set(map(lambda word: word.lower().strip(), row['question2'].split(" ")))
            return 1.0 * len(w1 & w2)
        df['word_Common'] = df.apply(normalized_word_Common, axis=1)

        def normalized_word_Total(row):
            w1 = set(map(lambda word: word.lower().strip(), row['question1'].split(" ")))
            w2 = set(map(lambda word: word.lower().strip(), row['question2'].split(" ")))
            return 1.0 * (len(w1) + len(w2))
        df['word_Total'] = df.apply(normalized_word_Total, axis=1)

        def normalized_word_share(row):
            w1 = set(map(lambda word: word.lower().strip(), row['question1'].split(" ")))
            w2 = set(map(lambda word: word.lower().strip(), row['question2'].split(" ")))
            return 1.0 * len(w1 & w2)/(len(w1) + len(w2))
        df['word_share'] = df.apply(normalized_word_share, axis=1)

        df['freq_q1+q2'] = df['freq_qid1']+df['freq_qid2']
        df['freq_q1-q2'] = abs(df['freq_qid1']-df['freq_qid2'])

        df.to_csv("df_fe_without_preprocessing_train.csv", index=False)

    # print the first 5 rows of the modified data frame:
    df.head()
```

Out[14]:

```

def get_token_features(q1,q2):
    """
    Function to get the first 10 Normal Features.
    """

    # Since we have 10 Normal Features:
    token_features = [0.0]*10

    # Converting the sentence into tokens:
    q1_tokens = q1.split()
    q2_tokens = q2.split()

    # If either q1 or q2 is empty, return token_features
    if len(q1_tokens) == 0 or len(q2_tokens) == 0:
        return token_features

    # Get the non-stopwords in q1 and q2:
    q1_words = set([word for word in q1_tokens if word not in STOP_WORDS])
    q2_words = set([word for word in q2_tokens if word not in STOP_WORDS])

    # Get the stopwords in q1 and q2:
    q1_stops = set([word for word in q1_tokens if word in STOP_WORDS])
    q2_stops = set([word for word in q2_tokens if word in STOP_WORDS])

    # Get the common non-stopwords from q1 & q2:
    common_word_count = len(q1_words.intersection(q2_words))
    # common_word_count = len(q1_words & q2_words)

    # Get the common stopwords from q1 & q2:
    common_stop_count = len(q1_stops.intersection(q2_stops))
    # common_stop_count = len(q1_stops & q2_stops)

    # Get the common tokens from q1 & q2:
    common_token_count = len(set(q1_tokens).intersection(set(q2_tokens)))
    # common_token_count = len(set(q1_tokens) & set(q2_tokens))

    # cwc_min:
    token_features[0] = common_word_count / (min(len(q1_words), len(q2_words)) + SAFE_DIV)
    # cwc_max:
    token_features[1] = common_word_count / (max(len(q1_words), len(q2_words)) + SAFE_DIV)

    # csc_min:
    token_features[2] = common_stop_count / (min(len(q1_stops), len(q2_stops)) + SAFE_DIV)
    # csc_max:
    token_features[3] = common_stop_count / (max(len(q1_stops), len(q2_stops)) + SAFE_DIV)

    # ctc_min:
    token_features[4] = common_token_count / (min(len(q1_tokens), len(q2_tokens)) + SAFE_DIV)
    # ctc_max:

```


0.358 0.358 0.357 0.358 0.357 0.358 0.356 0.358 0.359 0.359
0.357 0.356 0.357 0.357 0.357 0.356 0.358 0.359 0.358 0.36]

From the plot above, we can see that the lowest value of Misclassification error is generated in between $k=[20, 21, \dots, 40]$. That's the reason, we got our optimal_k to be 27.

Let us see the accuracy score after querying the k-NN model with the test data.

```
In [40]: # KNN with k = optimal_k
st = time()
# Configured parameters are:-
#
# 1. algorithm = 'auto':
#     automatically choose the algorithm (KDTree, BallTree or Brute Force)
#
# 2. metric = 'minkowski', p = 2:
#     Use L2 Minkowski Distance which is nothing but Euclidean Distance.
#
# 3. n_jobs = -1:
#     Use all the CPU cores to apply KNN Classification.

# Instantiate the learning model:
knn_optimal = KNeighborsClassifier(
    n_neighbors = optimal_k,
    algorithm = 'auto',
    metric = 'minkowski',
    p = 2,
    n_jobs = 3
)

# Fitting the model on train:
knn_optimal.fit(X_train, y_train)

# Predict the response on test:
predict_y_test = knn_optimal.predict(X_test)

# Evaluate the test accuracy:
acc_test = accuracy_score(predict_y_test, y_test, normalize=True) * float(100)
print('\n\nThe Accuracy of k-NN classifier on Quora Question Pairs Dataset
for predicting whether two given questions have the same intent or not with
k={} is {}'.format(optimal_k, acc_test))

time_taken(st)
```

The Accuracy of k-NN classifier on Quora Question Pairs Dataset
for predicting whether two given questions have the same intent or not with
k=27 is 65.70666666666666%

Runtime: 111.65 seconds

Conclusion

- Successfully applied k-Nearest Neighbours Algorithm on the Quora Question Pair Similarity Dataset as two variants:
 - k-NN after Simple Cross Validation [Train-Test Split: 70-30 of 25k Sampled Data Points]
 - k-NN after K-fold Cross Validation [Here, 'K' is not the same 'k' in k-NN].
- The Accuracy of k-NN classifier on a sample of 25,000 Data Points of Quora Question Pairs Dataset for predicting whether two given questions have the same intent or not with $k=27$ is 65.70666666666666%.
- Due to a computational discrepancy, only 25k Data Points were taken. But, if we apply the same algorithm on the entire Dataset, the Accuracy Score will be near to ~90%.

Thank You!

**Any
Questions/Queries?**
